



Training and Validation Loss Curves

— n=50 — n=200 — n=600 — n=1200


2stage reward fitting



FY end-to-end surrogate



SPO+ end-to-end surrogate



Dashed = train split; solid = validation split. Loss scales differ across panels.